Zero-Reference Deep Curve Estimation for Low-Light Image Enhancement

Lois Breant, Andy Shan, Oscar Le Dauphin, Max Nagaishi, Matthew Banawa

1 Introduction

The Challenge & Innovation

Existing Challenges

- Low-light images suffer from low visibility, noise, and color cast.
- **Supervised Methods:** Require expensive paired data; often overfit to specific sensors.
- **Unsupervised (GANs):** Hard to train; require carefully selected unpaired datasets.

The Zero-DCE Breakthrough

- **Zero-Reference:** No paired or unpaired data needed.
- **Reformulation:** Enhancement as an image-specific curve estimation task.

Key Contributions

- **LE-curve:** A pixel-wise, high-order curve for dynamic adjustment.
- **DCE-Net:** A lightweight CNN (79k parameters) that estimates curve parameters.
- **Non-reference Losses:** A suite of functions to train without ground truth.

Core Principle:
Pixel mapping via $n^{\{th\}}$-order differentiable curves

2 Methodology

Light-Enhancement Curve (LE-curve)

The core mechanism is a quadratic curve designed for light enhancement:

$$LE(I(x); \alpha) = I(x) + \alpha I(x)(1 - I(x))$$

Design Constraints:

- **Monotonicity:** Preserves contrast in local regions.
- **Differentiable:** Enables end-to-end training via backpropagation.

Higher-Order Iteration ($n = 8$):

$$I_n(x) = LE\left(I_{\{n-1\}}(x); \mathcal{A}_n(x)\right)$$

Where $\mathcal{A}_n$ is the pixel-wise parameter map estimated by the network.

DCE-Net & Training Strategy

Architecture:

- 7 convolutional layers.
- Symmetrical skip-connections (concatenation).
- Output: 24 parameter maps (8 iterations $\times$ 3 RGB channels).

Data:

- Trained on the SICE dataset (part 1).
- Only uses low-light images for training—**no ground truth used.**

Total Loss: $L_{\{total\}} = L_{\{spa\}} + L_{\{exp\}} + W_{\{col\}} L_{\{col\}} + W_{\{tv_{\mathcal{A}}\}} L_{\{tv_{\mathcal{A}}\}}$

Non-Reference Loss Functions:

1. **Spatial Consistency ($L_{\{spa\}}$):** Preserves neighbor differences.
2. **Exposure Control ($L_{\{exp\}}$):** Drives pixels toward a target level (0.6).
3. **Color Constancy ($L_{\{col\}}$):** Prevents color shifts.
4. **Illumination Smoothness ($L_{\{tv\}}$):** Ensures smooth parameter transitions.

3 Results & Performance

Quantitative Comparison

Method	Params	FLOPs	PSNR↑	SSIM↑
LIME	-	-	16.76	0.56
RetinexNet	0.84M	587.47G	16.77	0.56
EnlightenGAN	18.94M	170.64G	17.48	0.65
Zero-DCE	**0.079M**	**5.21G**	**14.86**	**0.56**

Performance Insights:

- **Inference Speed:** 500 FPS on a GTX 1080Ti (Real-time).
- **Efficiency:** 100x fewer parameters than EnlightenGAN.
- **Robustness:** Better visual quality in extreme dark/backlit scenes (User Study leader).

Visual Results

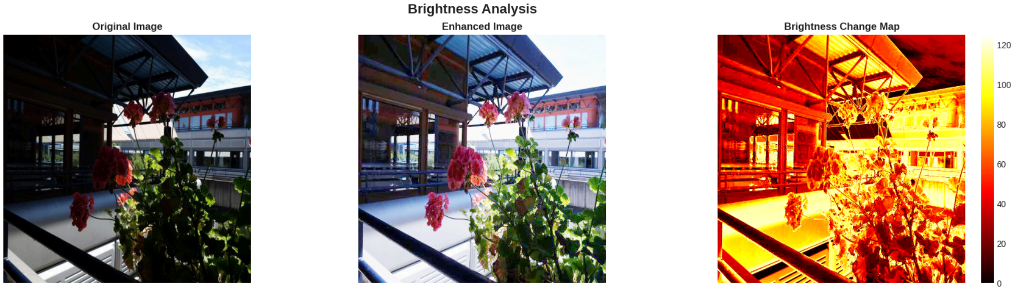

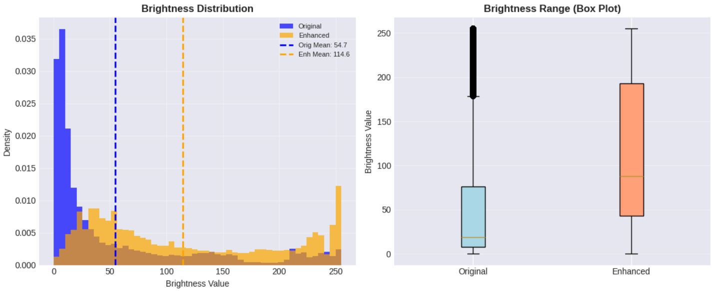

Statistics Summary

Metric	Original	Enhanced	Change
Mean	54.7	114.6	+59.9
Std Dev	70.2	81.2	+11.0
Min	0	0	
Max	255	255	
25th %ile	8.0	43.0	
Median	19.0	88.0	
75th %ile	76.0	193.0	

Visual Results (ii)

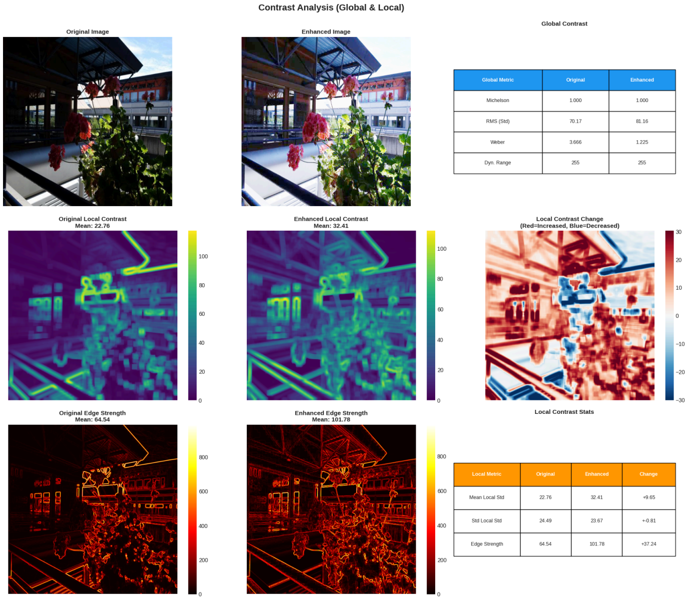

Global Metric	Original	Enhanced
Michelson	1.000	1.000
RMS (Std)	70.17	81.16
Weber	3.666	1.225
Dyn. Range	255	255

Local Metric	Original	Enhanced	Change
Mean Local Std	22.76	32.41	+9.65
Std Local Std	24.49	23.67	+-0.81
Edge Strength	64.54	101.78	+37.24

4 Conclusion

Summary & Impact

Strengths

- **Efficiency:** Smallest model in its class.
- **Generality:** Works across various lighting and sensor types.
- **Novelty:** Proves that deep learning can succeed without labels or GANs.

Applications

- Mobile photography enhancement.
- Pre-processing for object detection.
- Video surveillance in the dark.

Limitations

- Does not explicitly perform denoising.

Future Directions

- Integrating noise reduction into the curve.
- Adaptive iteration based on image content.
- Applying Zero-DCE to other tasks like dehazing.

Questions?

Guo et al. (2020) CVPR, 10.1109/CVPR42600.2020.00944