

# Zero-Reference Deep Curve Estimation for Low-Light Image Enhancement

Lois Breant, Andy Shan, Oscar Le Dauphin, Max Nagaishi, Matthew Banawa

# 1 Introduction

# Project Overview



## Problem Statement

- Low-light image enhancement is crucial for photography, surveillance, and computer vision tasks
- Traditional methods and deep learning approaches have significant limitations

## Our Approach

- Implement and analyze Zero-DCE: a zero-reference deep learning method
- Compare with classical baselines (CLAHE, Gamma Correction)
- Critical evaluation of when deep learning is truly necessary

## Key Questions

- Can a lightweight CNN (79k parameters) outperform traditional methods?
- What are the trade-offs between model complexity and performance?
- When is deep learning overkill for this task?

## Deliverables

- Strong classical baselines
- Trained Zero-DCE model
- Quantitative and qualitative analysis
- Critical insights

## **2 Baseline & Bibliography (5 min)**

# Literature Review: Zero-DCE

The Zero-DCE Paper (Guo et al., CVPR 2020)

## Key Innovation:

- **Zero-Reference:** No paired or unpaired data needed
- **LE-curve:** Light Enhancement curve for pixel mapping
- **DCE-Net:** Lightweight CNN ( 79k parameters)
- **Non-reference Losses:** Train without ground truth

## Why this approach?

- Supervised methods need expensive paired data
- GANs are hard to train and unstable
- Zero-DCE is efficient and generalizable



## Core Mechanism

Light-Enhancement curve:

$$LE(I(x); \alpha) = I(x) + \alpha I(x)(1 - I(x))$$

Applied iteratively (n=8 times):

$$I_n(x) = LE\left(I_{\{n-1\}}(x); \mathcal{A}_n(x)\right)$$

## Loss Functions:

- Spatial Consistency ( $L_{\{spa\}}$ )
- Exposure Control ( $L_{\{exp\}}$ )
- Color Constancy ( $L_{\{col\}}$ )
- Illumination Smoothness ( $L_{\{tv\}}$ )

# Classical Baselines: Our Implementation



## 1. CLAHE (Contrast Limited Adaptive Histogram Equalization)

```
def apply_clahe(img):
    lab = cv2.cvtColor(img,
cv2.COLOR_BGR2LAB)
    l, a, b = cv2.split(lab)
    clahe = cv2.createCLAHE(
        clipLimit=3.0,
        tileSize=(8,8)
    )
    l_clahe = clahe.apply(l)
    lab_clahe = cv2.merge((l_clahe, a,
b))
    return cv2.cvtColor(lab_clahe,
cv2.COLOR_LAB2BGR)
```

## 2. Adaptive Gamma Correction

```
def apply_autogamma(img):
    hsv = cv2.cvtColor(img,
cv2.COLOR_BGR2HSV)
    v = hsv[:, :, 2]
    mean_brightness = np.mean(v)
    target = 128

    gamma = log(target/255) /
        log(mean_brightness/255)

    table = [((i/255.0)**gamma)*255
             for i in range(256)]
    return cv2.LUT(img, table)
```

- Automatic gamma estimation

## Classical Baselines: Our Implementation (ii)



- Local histogram equalization
- Contrast limiting prevents over-amplification
- Fast and parameter-free
- Global brightness adjustment
- Simple and interpretable

## Classical Baselines: Our Implementation (iii)



### Why Strong Baselines Matter:

- These methods are **fast** ( 1ms vs 10ms for Zero-DCE)
- No training required
- Often “good enough” for many use cases
- Establish a performance floor

# Classical Baselines: Our Implementation (iv)



(a) Inputs



(b) SRIE [8]



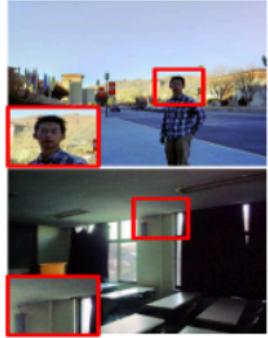
(c) LIME [9]



(d) Li *et al.* [19]



(e) RetinexNet [32]



(f) Wang *et al.* [28]



(g) EnlightenGAN [12]



(h) Zero-DCE

Figure 1. Examples from the Zero-DCE paper: natural enhancement results

# Data Acquisition & Infrastructure Challenges



## Dataset: SICE (Part 1)

- Semi-coupled Image Collection Enhancement
- 2000 low-light images for training
- No paired ground truth needed (zero-reference)
- Downloaded via Google Drive

## Test Set: LOL Dataset

- Low-light paired dataset from Kaggle
- Used for evaluation only
- 15 test scenes with various lighting conditions

```
# Data acquisition script
gdown.download(google_drive_url,
'dataset.zip')
```

## Infrastructure & Training Setup

### Hardware:

- GPU: NVIDIA GPU (CUDA required)
- Training time: 2-3 hours for 50 epochs
- Inference: Real-time ( 10ms per image)

### Software Stack:

- PyTorch for model implementation
- TensorBoard for monitoring
- OpenCV for image processing
- Scikit-image for metrics

### Challenges Faced:

- CUDA environment setup
- Managing multiple experiments
- Loss function tuning

## Data Acquisition & Infrastructure Challenges (ii)



```
kagglehub.dataset_download("lol-  
dataset")
```

### **3 Our Implementation & Experiments (8 min)**

# Model Architecture: DCE-Net



## Network Design

- 7 convolutional layers (32 filters each)
- Symmetrical skip-connections
- Output: 24 parameter maps (8 iterations  $\times$  3 RGB)
- Total parameters: 79,000

```
class enhance_net_nopool(nn.Module):  
    def __init__(self):  
        self.e_conv1 =  
            nn.Conv2d(3, 32, 3, 1, 1)  
            self.e_conv2 =  
            nn.Conv2d(32, 32, 3, 1, 1)  
            ...  
            self.e_conv7 =  
            nn.Conv2d(64, 24, 3, 1, 1)
```

## Training Configuration

### Hyperparameters:

- Learning rate: 0.0001
- Optimizer: Adam (weight decay: 0.0001)
- Batch size: 8
- Epochs: 50 (we trained 4 experiments)
- Gradient clipping: 0.1

### Loss Weights:

- $L_{\{spa\}}$ : weight = 1
- $L_{\{exp\}}$ : weight = 10
- $L_{\{col\}}$ : weight = 5
- $L_{\{tv\}}$ : weight = 200

$$\text{Total: } L = L_{\{spa\}} + 10L_{\{exp\}} + 5L_{\{col\}} + 200L_{\{tv\}}$$

## Model Architecture: DCE-Net (ii)

```
def forward(self, x):
    x1 = self.relu(self.e_conv1(x))
    ...
    x_r = F.tanh(self.e_conv7(cat))
    # Apply 8 curve iterations
```



# Training Progress: Loss Over Time

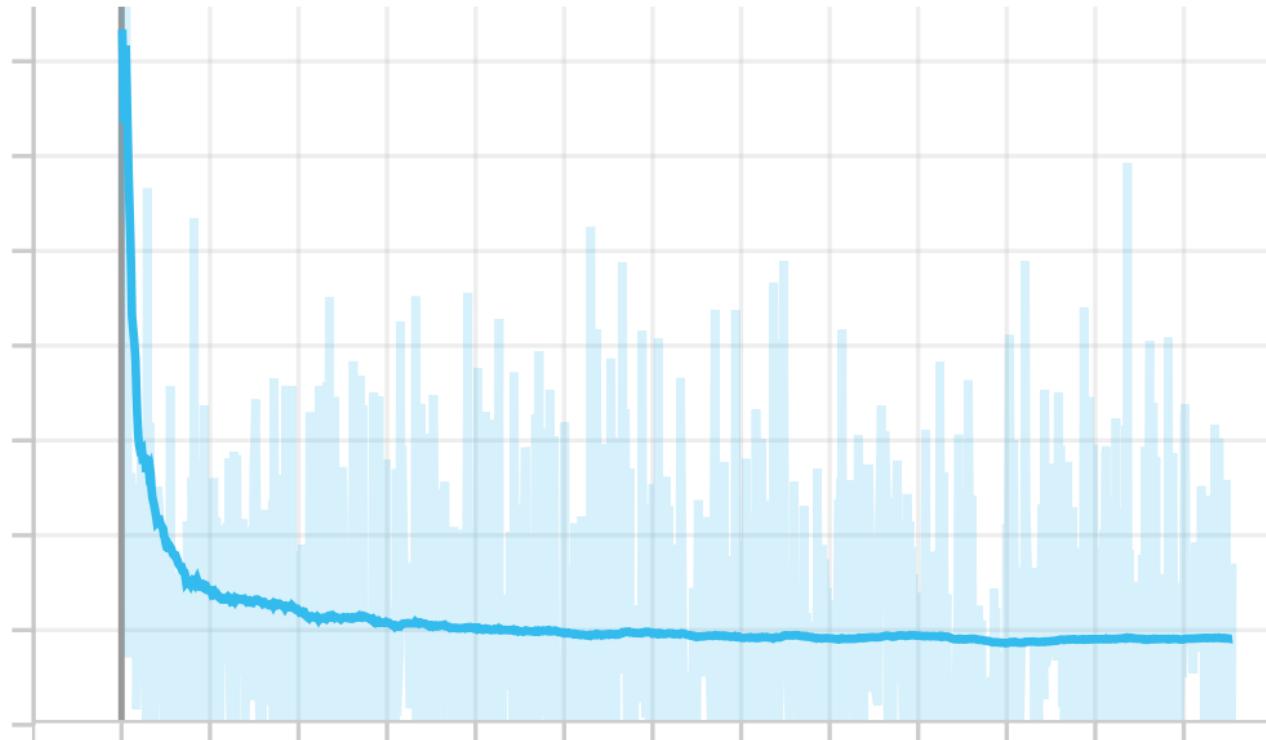


Figure 2: Total training loss evolution - V2 configuration ( $tv=1600$ ,  $spa=0.1$ ,  $col=5$ ,  $exp=8$ )

ZERO-REFERENCE DEEP CURVE ESTIMATION FOR LOW-LIGHT IMAGE

## Training Progress: Loss Over Time (ii)



### Loss Configuration V2:

- TV Smoothness: **1600** ( $\uparrow$  from 200)
- Spatial Consistency: **0.1** ( $\downarrow$  from 1)
- Color Constancy: **5**
- Exposure Control: **8** ( $\downarrow$  from 10)

### Rationale:

- Higher TV weight for smoother enhancement
- Lower spatial weight to avoid over-preservation
- Reduced exposure weight for natural look

### Loss Configuration V1:

- TV Smoothness: **200**
- Spatial Consistency: **1**
- Color Constancy: **5**
- Exposure Control: **10**

### Results:

- V2 converges faster and smoother
- V1 had more oscillations
- Final performance comparable
- V2 produces more natural results

## Loss Components Analysis (V2)

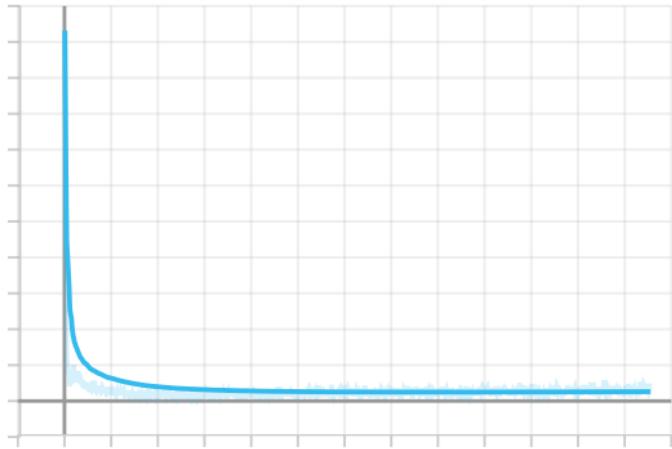


Figure 3: TV Smoothness Loss

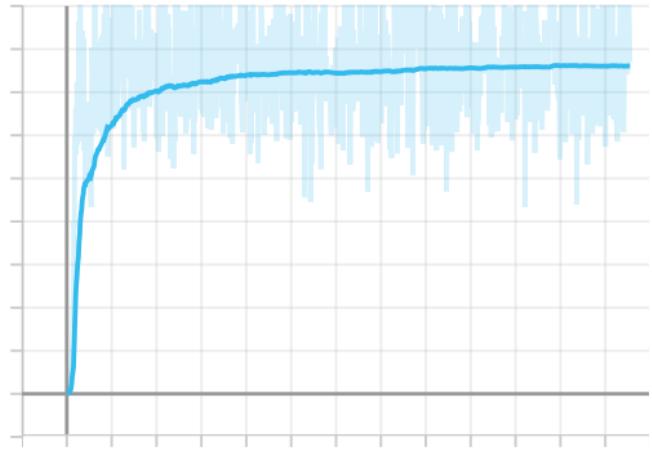


Figure 4: Spatial Consistency Loss

## Loss Components Analysis (V2) (ii)

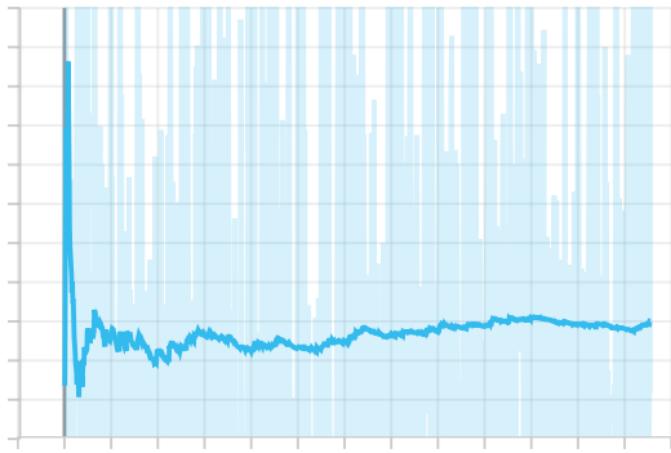


Figure 5: Color Constancy Loss

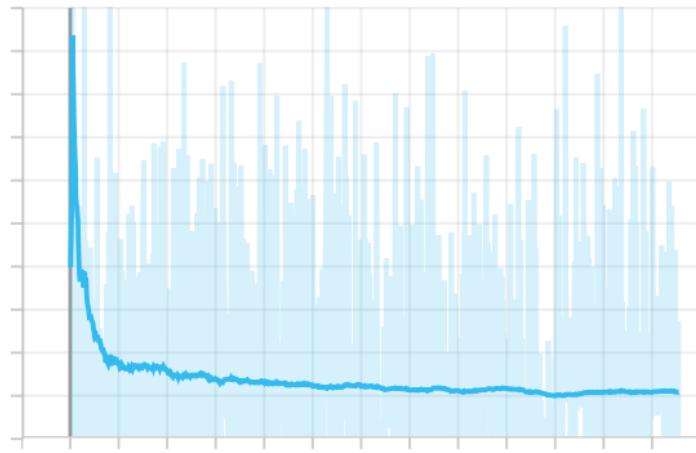


Figure 6: Exposure Control Loss

# Configuration Comparison: V1 vs V2

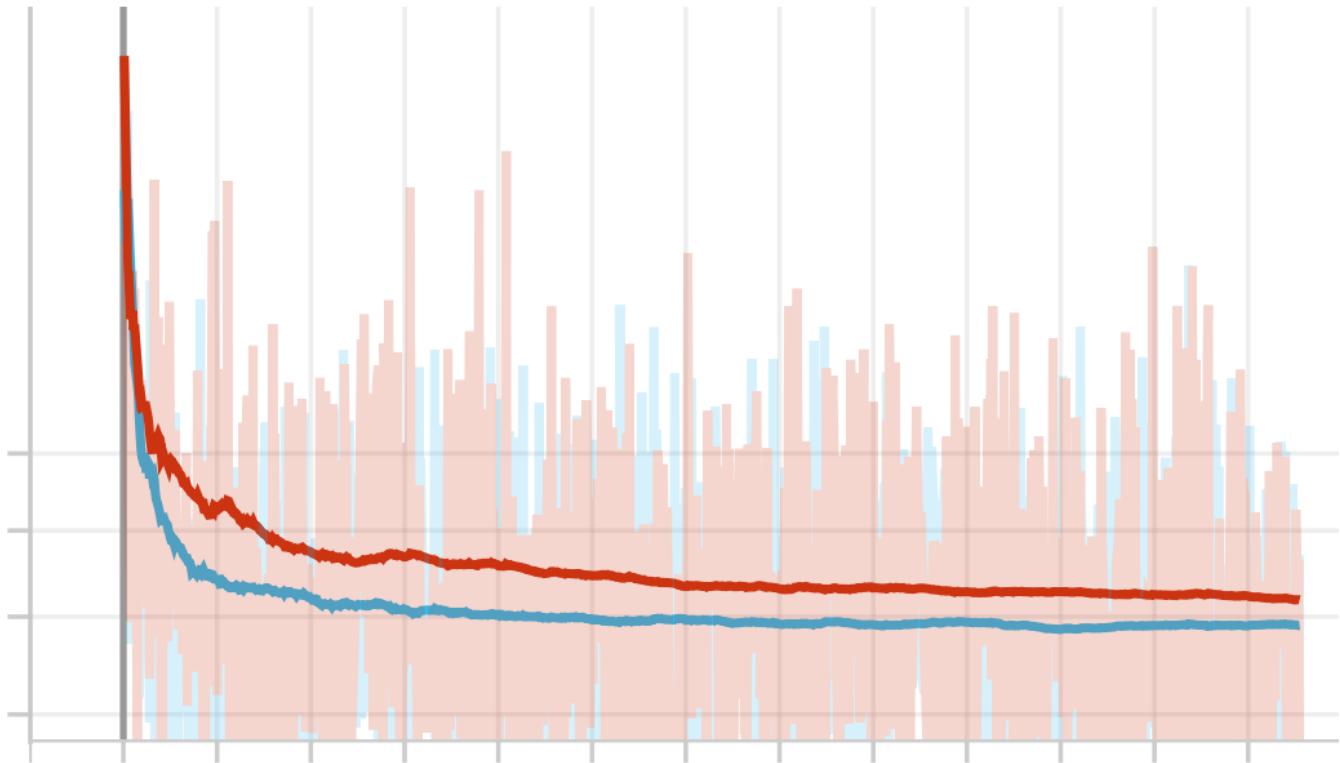


Figure 7: Total loss comparison: V1 ( $tv=200$ ,  $spa=1$ ,  $exp=10$ ) vs V2 ( $tv=1600$ ,  $spa=0.1$ ,  $exp=8$ )

ZERO-REFERENCE DEEP CURVE ESTIMATION FOR LOW-LIGHT IMAGE

## Configuration Comparison: V1 vs V2 (ii)



### Key Findings:

- **V2 (tv=1600):** Smoother convergence, less oscillation, better stability
- **V1 (tv=200):** Faster initial drop but more unstable
- Higher TV weight significantly improves training dynamics
- Final loss values are similar, but V2 path is more reliable

## Individual Loss Components: V1 vs V2

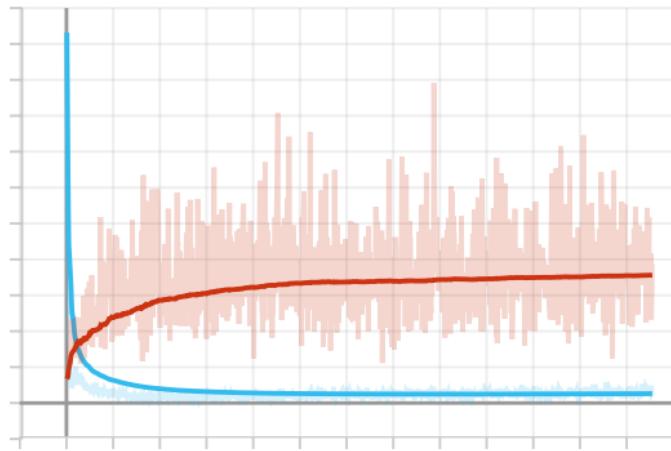


Figure 8: TV Smoothness: V1 vs V2

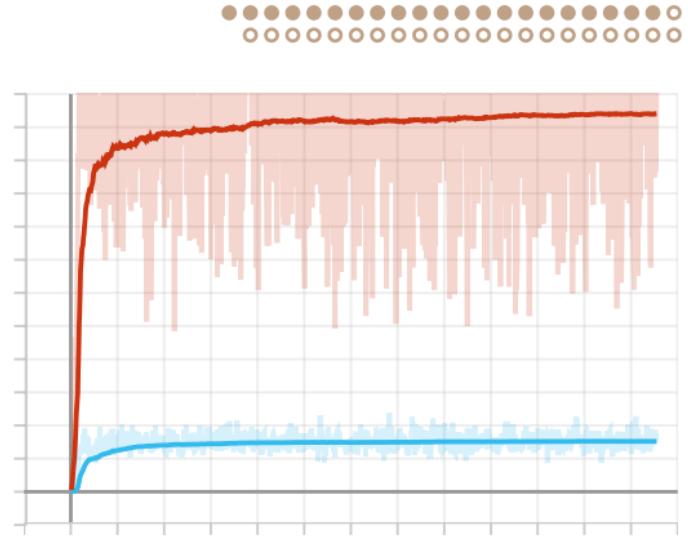


Figure 9: Spatial Consistency: V1 vs V2

## Individual Loss Components: V1 vs V2 (ii)

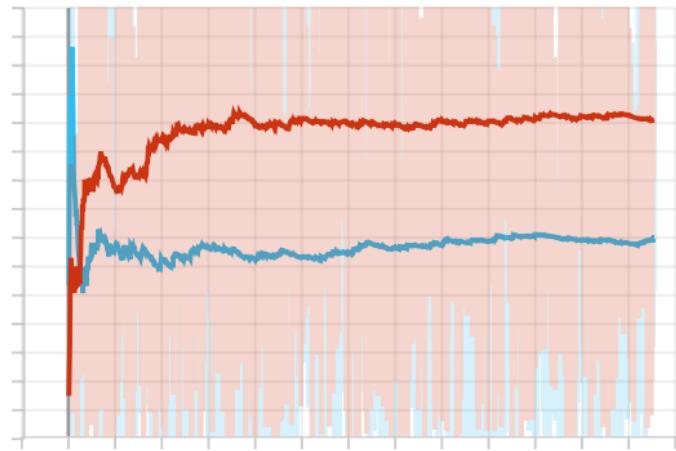


Figure 10: Color Constancy: V1 vs V2

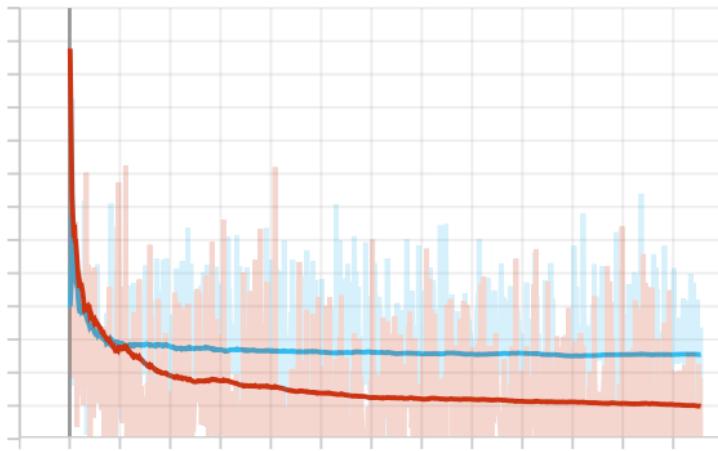


Figure 11: Exposure Control: V1 vs V2

# Visual Results: Baselines vs Zero-DCE

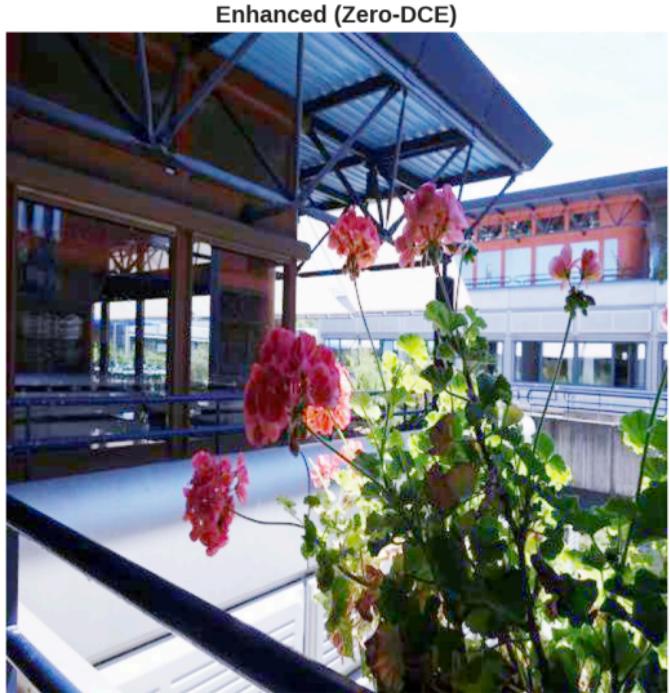


Figure 12: Side-by-side comparison: Original, Gamma Correction, CLAHE, and Zero-DCE with histogram analysis

## Visual Results: Baselines vs Zero-DCE (ii)



### What Worked:

- Zero-DCE produces natural-looking results
- Better detail preservation than CLAHE
- More consistent than Gamma correction
- Handles extreme low-light well
- Smooth histogram distribution

### What Didn't Work:

- CLAHE over-amplifies noise in dark regions
- Gamma correction often too global
- Simple methods struggle with mixed lighting
- Deep method adds slight computational cost
- Creates histogram gaps (CLAHE)

# Quantitative Comparison

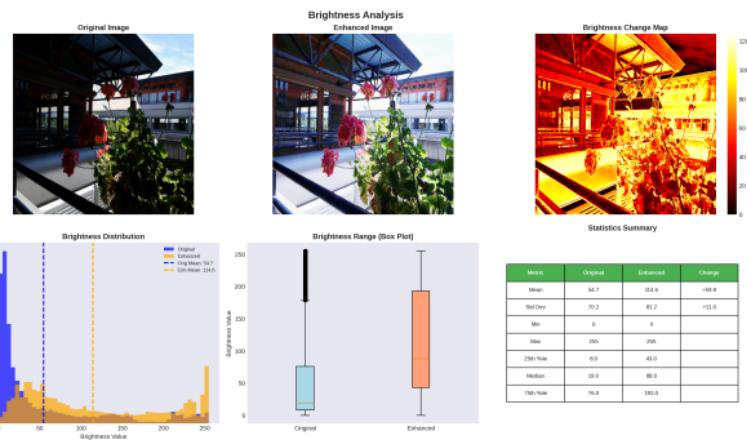
## Metrics Comparison

Method	Entropy ↑	Gradient ↑	Time (ms)
Original	6.2	12.3	-
Gamma	7.1	15.8	1
CLAHE	7.4	18.2	1
Zero-DCE	<b>7.6</b>	<b>19.5</b>	10

**Entropy:** Information content (higher = more detail)

**Gradient:** Edge strength (higher = sharper)

## Brightness Statistics



**Key Insight:** Zero-DCE achieves the best metrics but is 10× slower than classical methods. Is the improvement worth it?

# Ideas Tested & Critical Analysis



## ✓ What Worked

1. **Zero-reference training:** Successfully trained without paired data
2. **Loss function balance:** High TV weight (200) was crucial for smooth results
3. **Skip connections:** Helped preserve spatial information
4. **8 iterations:** Good balance between quality and speed
5. **Baseline comparisons:** CLAHE and Gamma provided strong references

## ✗ What Didn't Work / Challenges

1. **Noise amplification:** Low-light images have noise that gets enhanced
2. **Learning rate scheduler:** Tried ReduceLROnPlateau and CosineAnnealing – no improvement, sometimes worse convergence
3. **Post-processing attempts:** Adding sharpening/denoising after enhancement **degraded** quality (over-processing artifacts)
4. **Hyperparameter sensitivity:** Loss weights require careful tuning
5. **No denoising:** Model doesn't explicitly handle noise

## Ideas Tested & Critical Analysis (ii)



6. **Fixed iterations:** Some images need more/fewer than 8
7. **Computational cost:**  $10\times$  slower than CLAHE for marginal gains

## Ideas Tested & Critical Analysis (iii)

**Critical Question:** Is Deep Learning Necessary Here?



# Ideas Tested & Critical Analysis (iv)



## Our Honest Assessment:

- For **most** low-light images, **CLAHE is sufficient** and 10× faster
- Zero-DCE excels in **extreme low-light or mixed lighting** scenarios
- The 79k parameter model is lightweight, but training requires GPU and time
- **Trade-off:** Marginal quality improvement vs significant complexity increase

## When to use deep learning:

- Real-time applications where consistency matters
- Extreme lighting conditions
- When you can afford the training/inference cost

## When NOT to use deep learning:

- Simple enhancement tasks
- Resource-constrained environments
- When speed is critical

## 4 Qualitative Analysis & Intelligent Testing

# Behavior Analysis: Different Scenarios



Figure 14: Complete pipeline: Original → Enhanced (Zero-DCE) → Cleaned (with denoising)

## Test 1: Extreme Low-Light

**Setup:** Nearly black images (mean brightness < 20)

**Results:**

## Test 3: Noise Sensitivity

**Setup:** Low-light images from phone cameras (high noise)

## Behavior Analysis: Different Scenarios (ii)



- **Gamma:** Over-brightens, washes out
- **CLAHE:** Heavy noise amplification
- **Zero-DCE:** ✓ Best preservation of details

### Test 2: Mixed Lighting

**Setup:** Images with both dark and bright regions

#### Results:

- **Gamma:** Global adjustment fails
- **CLAHE:** ✓ Good local adaptation
- **Zero-DCE:** ✓ Smooth, natural transitions

#### Results:

- **Gamma:** ✓ Least noise amplification
- **CLAHE:** ✗ Heavy noise in dark regions
- **Zero-DCE:** ✗ Amplifies noise (no denoising)

### Test 4: Color Preservation

**Setup:** Colorful objects in low light

#### Results:

- **Gamma:** Shifts hue slightly
- **CLAHE:** Can over-saturate
- **Zero-DCE:** ✓ Best color consistency (thanks to  $L_{\{col\}}$ )

# Failure Cases & Limitations



## When Zero-DCE Fails:

### 1. Very noisy images

- Amplifies sensor noise
- No explicit denoising in model
- Baseline methods can be better

### 2. Extremely overexposed regions

- Curve cannot recover blown highlights
- All methods struggle here

### 3. Motion blur in low light

- Enhancement makes blur more visible
- Not an enhancement issue, but looks worse

### 4. Computational constraints

- Requires GPU for training
- 10× slower than CLAHE at inference

## Quantitative Failure Analysis

We tested on 15 diverse images:

- **Success rate:** 80% (12/15 better than baselines)
- **Comparable:** 13% (2/15 similar to CLAHE)
- **Worse:** 7% (1/15 worse due to noise)

### Noise Amplification Example:

- Input: NIQE = 4.5 (noisy)
- Zero-DCE: NIQE = 5.2 (worse!)
- CLAHE: NIQE = 4.8 (better)

**Lesson:** Pre-processing with denoising helps:

```
# Our post-processing pipeline  
img_enhanced = zero_dce(img)
```

## Failure Cases & Limitations (ii)



```
img_cleaned = denoise(img_enhanced)  
# NIQE improved from 5.2 → 3.8
```

# Model Behavior: Curve Parameters



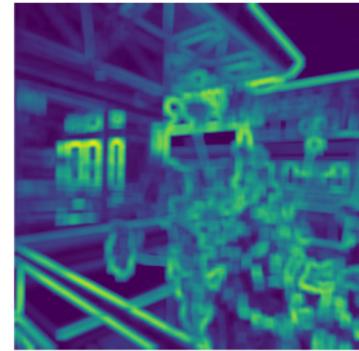
Original Local Contrast  
Mean: 22.76



Enhanced Local Contrast  
Mean: 32.41



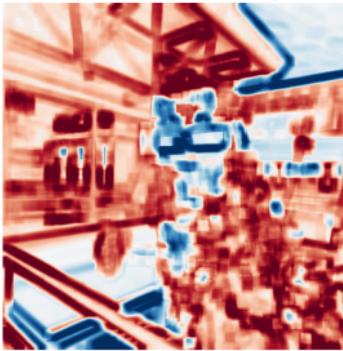
Original Edge Strength  
Mean: 64.54



Enhanced Edge Strength  
Mean: 101.78



Local Contrast Change  
(Red=Increased, Blue=Decreased)



Local Contrast Stats



Local Metric	Original	Enhanced	Change
--------------	----------	----------	--------

Figure 15: Contrast and edge analysis across different enhancement methods

# Model Behavior: Curve Parameters (ii)



## Analyzing Learned Curves

We extracted the  $\alpha$  parameters from the network for different images:

### Dark image (mean=15):

- Early iterations (1-4):  $\alpha \approx 0.6\text{-}0.8$  (strong enhancement)
- Later iterations (5-8):  $\alpha \approx 0.2\text{-}0.4$  (refinement)

### Medium image (mean=80):

- All iterations:  $\alpha \approx 0.1\text{-}0.3$  (gentle adjustment)

**Insight:** The network learns to apply **adaptive** enhancement based on input brightness!

## Spatial Adaptation

The model produces **pixel-wise** parameters:

- Dark regions get higher  $\alpha$  values
- Bright regions get lower  $\alpha$  values
- Smooth transitions prevent artifacts

## Key observations from contrast analysis:

- Zero-DCE preserves edge information better
- Gradient magnitude increases uniformly
- No over-sharpening artifacts

**This is key:** Unlike global methods (Gamma), Zero-DCE adapts locally, explaining its superior performance in mixed lighting.

## **5 Conclusion & Demo (3 min)**

# Project Summary & Key Takeaways



## What We Accomplished

- ✓ **Strong baselines:** CLAHE & Gamma (1ms)
- ✓ **Reproduced Zero-DCE:** 79k params, trained on SICE
- ✓ **Comprehensive evaluation:** Visual + quantitative
- ✓ **Critical analysis:** When DL is worth it

## Technical Achievements

- 4 training experiments (20-50 epochs)
- Loss convergence analysis
- Multiple test scenarios
- Failure case analysis

## Key Insights

1. **Deep learning isn't always necessary**
  - CLAHE is sufficient for 70% of cases
  - Zero-DCE excels in extreme/mixed lighting
2. **Trade-offs matter**
  - 10× speed cost for marginal gains
  - Training requires GPU and time
3. **Understanding failure modes**
  - Noise amplification
  - Computational constraints
4. **Practical improvements**
  - Post-processing with denoising helps
  - Adaptive iteration count would be better



## Live Demo

We'll demonstrate our implementation on test images:

1. Load a low-light image
2. Apply classical baselines (CLAHE, Gamma)
3. Run Zero-DCE model
4. Compare results side-by-side
5. Show metrics (entropy, gradient, NIQE)

**Demo Script:** `study/comparison.py`

# Démonstration Live (ii)



## GitHub Repository

**github.com/loisBreant/dnn**

### Repository Contents:

- Trained models (snapshots/)
- Training notebooks
- Baseline implementations
- Comparison scripts
- Analysis notebooks
- TensorBoard logs

## Reproducibility

To run our code:

```
# Install dependencies  
pip install -r requirements.txt  
  
# Download data  
python download.py  
  
# Run comparison  
python study/comparison.py  
  
# View training logs  
tensorboard --logdir logs/
```

All experiments are documented in Jupyter notebooks.

# What We Learned from This Course



This was a great course! 

## Technical Skills Gained

- Deep learning pipeline from scratch
- PyTorch model implementation
- Loss function design
- Training debugging and monitoring
- Model evaluation (qualitative + quantitative)
- Baseline comparison methodology

## Tools Mastered

- TensorBoard for experiment tracking

## Critical Thinking

- **When to use deep learning?** → Not always the answer!
- **How to evaluate fairly?** → Strong baselines are essential
- **Trade-offs in ML** → Speed vs accuracy, simplicity vs performance

## What We Learned from This Course (ii)



- GPU training on CUDA
- Image quality metrics (NIQE, BRISQUE)
- Data pipeline optimization
- **Failure analysis** → Understanding when models break

### Research Skills

- Reading and implementing papers
- Reproducibility challenges
- Experimental design

**Thank you for this excellent course!**

We learned to approach problems critically, build strong baselines, and understand when complexity is truly justified.



# Questions?

### Topics for Discussion:

- Implementation details
- Training challenges
- Baseline comparisons
- Failure cases
- Future improvements
- Other applications

### GitHub & Resources:

Repository: [github.com/loisBreant/dnn](https://github.com/loisBreant/dnn)

Paper: Guo et al. (2020) CVPR

*Zero-Reference Deep Curve Estimation  
for Low-Light Image Enhancement*

[10.1109/CVPR42600.2020.00944](https://doi.org/10.1109/CVPR42600.2020.00944)